

EECS E6895 Advanced Big Data Analytics

Homework 3 (using grace day)

Name: Jingyi Yuan

UNI: jy2736

Algorithm 1: Matrix multiplication

I use CUDA to calculate the product of two matrices. The first matrix is an A*B matrix and the second one is a B*C one.

I write both the CPU and the GPU way to compute the product and record the start time and end time of these two ways.

CPU:

```
clock_t startc, finishc;

//start using CPU to multiply matrix
startc = clock();
for (int i = 0; i < A; i++){
    for (int j = 0; j < C; j++){
        {
            float sum = 0;
            for (int k = 0; k < B; k++){
                {
                    sum += a[i * B + k] * b[k * C + j];
                }
            }
            c[i * C + j] = sum;
        }
    }
    finishc = clock();
```

GPU:

```
#define A 2
#define B 3
#define C 4

__global__ void product(float *a, float *b, float *c, int aa, int bb, int cc)
{
    int ix = threadIdx.x + blockIdx.x * blockDim.x;
    int iy = threadIdx.y + blockIdx.y * blockDim.y;
    if (ix < aa && iy < cc)
    {
        float sum = 0;
        for (int index = 0; index < bb; index++){
            {
                sum += a[ix * bb + index] * b[index * cc + iy];
            }
        }
        c[ix * cc + iy] = sum;
    }
}
```

```

//***** GPU, matrix initialization *****
clock_t start, finish;
float *d_a, *d_b, *d_c;

//start using GPU to multiply matrix
start = clock();

cudaMalloc(&d_a, A * B*sizeof(float));
cudaMalloc(&d_b, B * C*sizeof(float));
cudaMalloc(&d_c, A * C*sizeof(float));

cudaMemcpy(d_a, a, A * B*sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, B * C*sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_c, c, A * C*sizeof(float), cudaMemcpyHostToDevice);

int dimx = 32;
int dimy = 32;
dim3 block(dimx, dimy);
dim3 grid((A + block.x - 1) / block.x, (C + block.y - 1) / block.y);

product << <grid, block >> >(d_a, d_b, d_c, A, B, C);

cudaMemcpy(c, d_c, A * C*sizeof(float), cudaMemcpyDeviceToHost);

```

I first define A = 2, B = 3, C = 4 to verify the correctness of the code and the output is:

```

input matrix 1:
2.00    2.00    2.00
2.00    2.00    2.00

input matrix 2:
2.00    2.00    2.00    2.00
2.00    2.00    2.00    2.00
2.00    2.00    2.00    2.00

output matrix (CPU):
12.00    12.00    12.00    12.00
12.00    12.00    12.00    12.00
*****
The total time using CPU: 0.000000 seconds
*****

output matrix (using GPU):
12.00    12.00    12.00    12.00
12.00    12.00    12.00    12.00
*****
The total time using GPU: 0.171000 seconds
*****

```

We can see that the result is correct, and when matrix is small, CPU is faster than GPU since we need to initialize when using GPU.

Then I used large matrix to do the calculation and change the value of A, B and C.

```
#define A 1000
#define B 1500
#define C 2000
```

This time, I only outputted one value in the product since I initialized every number in input1 and input 2 to be 2, thus every number in the output is the same and I only need to ensure that the number is within the range of float and does not overflow.

```
number in matrix: 6000.00
*****
The total time using CPU: 23.146000 seconds
*****
*****
The total time using GPU: 1.612000 seconds
*****
```

We can see that this time GPU is much faster than CPU. When dealing with large dataset, GPU parallel computing largely improves the speed of the program.

Algorithm 2: Linear Regression

In the hw2, I used linear regression to predict stock price using the returns of 7 days. In this homework, I applied linear regression to a simple dataset x and y: x = [0 1 2 3 4 5] and y=[0 20 60 68 77 110].

To do linear regression, I use ordinary least square to calculate β . To do this, I set the first column of input X to be 1 and the second column to be x. That is:

$$\mathbf{Y} = \begin{bmatrix} Y_1 \\ \vdots \\ Y_n \end{bmatrix} \quad \mathbf{X} = \begin{bmatrix} 1 & X_{11} & \dots & X_{1p} \\ \vdots & & & \\ 1 & X_{n1} & \dots & X_{np} \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_0 \\ \vdots \\ w_p \end{bmatrix}$$

$$\mathbf{Y} = \mathbf{X}\mathbf{w} + \epsilon$$

$$Y_i = w_0 + w_1X_{i1} + \dots + w_pX_{ip}$$

(reference: STAT W4240 Section 1 Data Mining Giovanni Motta Lecture9_Sec01)

And X and y look like:

```

X:
1.0000  0.0000
1.0000  1.0000
1.0000  2.0000
1.0000  3.0000
1.0000  4.0000
1.0000  5.0000

y:
0.0000  20.0000  60.0000  68.0000  77.0000  110.0000

```

I use the equation:

$$\hat{\mathbf{w}} = \left(\mathbf{X}^T \mathbf{X} \right)^{-1} \mathbf{X}^T \mathbf{Y}$$

to calculate w.

First, calculate X transpose:

```

#define A 6
#define B 2
#define C 1

__global__ void transpose(float *odata, float* idata, int ny, int nx)
{
    int ix = blockDim.x * blockIdx.x + threadIdx.x;
    int iy = blockDim.y * blockIdx.y + threadIdx.y;

    if (ix < nx && iy < ny)
    {
        odata[ix * ny + iy] = idata[iy * nx + ix];
    }
}

//----- matrix initialization for transpose -----
float *d_a, *d_b;

cudaMalloc(&d_a, A * B*sizeof(float));
cudaMalloc(&d_b, B * A*sizeof(float));

cudaMemcpy(d_a, X, A * B*sizeof(float), cudaMemcpyHostToDevice);
cudaMemcpy(d_b, y, B * A*sizeof(float), cudaMemcpyHostToDevice);

int dimx = 32;
int dimy = 32;
dim3 block(dimx, dimy);
dim3 grid((B + block.x - 1) / block.x, (A + block.y - 1) / block.y);

transpose << <grid, block >> >(d_b, d_a, A, B);

cudaMemcpy(b, d_b, B * A*sizeof(float), cudaMemcpyDeviceToHost);

```

Then calculate the product of X transpose and X:

```

__global__ void product(float *a, float *b, float *c, int aa, int bb, int cc)
{
    int ix = blockIdx.x * blockDim.x + threadIdx.x;
    int iy = blockIdx.y * blockDim.y + threadIdx.y;
    if (ix < aa && iy < cc)
    {
        float sum = 0;
        for (int index = 0; index < bb; index++)
        {
            sum += a[ix * bb + index] * b[index * cc + iy];
        }
        c[ix * cc + iy] = sum;
    }
}

```

```

//----- matrix initialization for multiplication X' * X -----
float *c;
c = (float *)malloc(B * B*sizeof(float));
float *d_c;

//start using GPU to multiply matrix
cudaMalloc(&d_c, B * B*sizeof(float));
cudaMemcpy(d_c, c, B * B*sizeof(float), cudaMemcpyHostToDevice);

dim3 grid2((B + block.x - 1) / block.x, (B + block.y - 1) / block.y);

product << grid2, block >> (d_b, d_a, d_c, B, A, B);

cudaMemcpy(c, d_c, B * B*sizeof(float), cudaMemcpyDeviceToHost);

```

Then calculate the inverse of the product of X transpose and X:

```

#define PERR(call) \
if (call) {\
    fprintf(stderr, "%s:%d Error [%s] on \"%#call\"\n", __FILE__, __LINE__, \
        cudaGetErrorString(cudaGetLastError())); \
    exit(1); \
}

#define ERRCHECK \
if (cudaPeekAtLastError()) {\
    fprintf(stderr, "%s:%d Error [%s]\n", __FILE__, __LINE__, \
        cudaGetErrorString(cudaGetLastError())); \
    exit(1); \
}

```

```

__global__ void inv_kernel(float *a_i, float *c_o, int n)
{
    int *p = (int *)malloc(3 * sizeof(int));
    int *info = (int *)malloc(sizeof(int));
    int batch;
    cublasHandle_t hdl;
    cublasStatus_t status = cublasCreate_v2(&hdl);

    info[0] = 0;
    batch = 1;
    float **a = (float **)malloc(sizeof(float *));
    *a = a_i;
    const float **aconst = (const float **)a;
    float **c = (float **)malloc(sizeof(float *));
    *c = c_o;
    status = cublasSgetrfBatched(hdl, n, a, n, p, info, batch);
    __syncthreads();
    status = cublasSgetriBatched(hdl, n, aconst, n, p,
        c, n, info, batch);
    __syncthreads();
    cublasDestroy_v2(hdl);
}

static void run_inv(float *in, float *out, int n)
{
    float *a_d, *c_d;

    PERR(cudaMalloc(&a_d, n*n*sizeof(float)));
    PERR(cudaMalloc(&c_d, n*n*sizeof(float)));
    PERR(cudaMemcpy(a_d, in, n*n*sizeof(float), cudaMemcpyHostToDevice));

    inv_kernel << 1, 1 >> >(a_d, c_d, n);

    cudaDeviceSynchronize();
    PERRCHECK;

    PERR(cudaMemcpy(out, c_d, n*n*sizeof(float), cudaMemcpyDeviceToHost));
    PERR(cudaFree(a_d));
    PERR(cudaFree(c_d));
}

//----- pinv(X' * X) -----
float *invmatrix;
invmatrix = (float *)malloc(B * B*sizeof(float));
run_inv(c, invmatrix, B);

```

Finally, use the product function to calculate the final result:

```
//----- pinv(X'*X)*X' -----
float *invma;
cudaMalloc(&invma, B * B*sizeof(float));
cudaMemcpy(invma, invmatrix, B * B*sizeof(float), cudaMemcpyHostToDevice);

float *e;
e = (float *)malloc(B * A*sizeof(float));
float *d_e;
cudaMalloc(&d_e, B * A*sizeof(float));
cudaMemcpy(d_e, e, B * A*sizeof(float), cudaMemcpyHostToDevice);

dim3 grid3((B + block.x - 1) / block.x, (A + block.y - 1) / block.y);
product << <grid3, block >> >(invma, d_b, d_e, B, B, A);
cudaMemcpy(e, d_e, B * A*sizeof(float), cudaMemcpyDeviceToHost);

//----- pinv(X'*X)*X'*y' -----
float *res;
res = (float *)malloc(B * C*sizeof(float));
float *d_res;
cudaMalloc(&d_res, B * C*sizeof(float));
cudaMemcpy(d_res, res, B * C*sizeof(float), cudaMemcpyHostToDevice);

float *d_y;
cudaMalloc(&d_y, A * C*sizeof(float));
cudaMemcpy(d_y, y, A * C*sizeof(float), cudaMemcpyHostToDevice);

dim3 grid4((B + block.x - 1) / block.x, (C + block.y - 1) / block.y);
product << <grid4, block >> >(d_e, d_y, d_res, B, A, C);
cudaMemcpy(res, d_res, B * C*sizeof(float), cudaMemcpyDeviceToHost);
```

Then we have:

```
the result is:
3.7619
20.8286

-----
The regression model is: y = 20.8286 * x + 3.7619
-----
```

The output of this part is:

```
C:\Users\hpan4\Desktop\lr\lr>cuda-memcheck ./kernel
===== CUDA-MEMCHECK
X:
1.0000  0.0000
1.0000  1.0000
1.0000  2.0000
1.0000  3.0000
1.0000  4.0000
1.0000  5.0000

y:
0.0000  20.0000  60.0000  68.0000  77.0000  110.0000

X':
1.0000  1.0000  1.0000  1.0000  1.0000  1.0000
0.0000  1.0000  2.0000  3.0000  4.0000  5.0000

X'*X:
6.0000  15.0000
15.0000  55.0000

pinv(X'*X):
0.5238, -0.1429,
-0.1429, 0.0571,

pinv(X'*X)*X:
0.5238  0.3810  0.2381  0.0952  -0.0476 -0.1905
-0.1429 -0.0857 -0.0286 0.0286  0.0857  0.1429

the result is:
3.7619
20.8286

-----
The regression model is:  $y = 20.8286 * x + 3.7619$ 
-----
===== ERROR SUMMARY: 0 errors
```


I went through this process in MATLAB and got the same result:

```
>> X'*X
ans =
     6     15
    15     55

>> pinv(X'*X)
ans =
     0.5238    -0.1429
    -0.1429     0.0571

>> pinv(X'*X)*X'
ans =
     0.5238     0.3810     0.2381     0.0952    -0.0476    -0.1905
    -0.1429    -0.0857    -0.0286     0.0286     0.0857     0.1429

>> pinv(X'*X)*X'*y
ans =
     3.7619
    20.8286
```

To verify this result, I use function `polyfit` in MATLAB and got exactly the same answer (`polyfit` change the position of $\beta(0)$ and $\beta(1)$).

```
>> x=[0 1 2 3 4 5];
>> y=[0 20 60 68 77 110];
>> coef=polyfit(x,y,1);
>> coef
```

```
coef =
    20.8286     3.7619
```

Thus this function can be used to do linear regression. And we can also apply this algorithm to large dataset.